

| Algorithm | Pros | Cons | Good at |
|--------------------------------|---|--|---|
| Linear regression | <ul style="list-style-type: none"> - Very fast (runs in constant time) - Easy to understand the model - Less prone to overfitting | <ul style="list-style-type: none"> - Unable to model complex relationships - Unable to capture nonlinear relationships without first transforming the inputs | <ul style="list-style-type: none"> - The first look at a dataset - Numerical data with lots of features |
| Decision trees | <ul style="list-style-type: none"> - Fast - Robust to noise and missing values - Accurate | <ul style="list-style-type: none"> - Complex trees are hard to interpret - Duplication within the same sub-tree is possible | <ul style="list-style-type: none"> - Spam classification - Medical diagnosis - Credit risk analysis |
| Neural networks | <ul style="list-style-type: none"> - Extremely powerful - Can model even very complex relationships - No need to understand the underlying data - Almost works by “magic” | <ul style="list-style-type: none"> - Prone to overfitting - Long training time - Requires significant computing power for large datasets - Model is essentially unreadable | <ul style="list-style-type: none"> - Images - Video - “Human-intelligence” type tasks like driving or flying - Robotics |
| Support Vector Machines | <ul style="list-style-type: none"> - Can model complex, nonlinear relationships - Robust to noise (because they maximize margins) | <ul style="list-style-type: none"> - Need to select a good kernel function - Model parameters are difficult to interpret - Sometimes numerical stability problems - Requires significant memory and processing power | <ul style="list-style-type: none"> - Classifying proteins - Text classification - Image classification - Handwriting recognition |
| K-Nearest Neighbors | <ul style="list-style-type: none"> - Simple - Powerful - No training involved (“lazy”) - Naturally handles multiclass classification and regression | <ul style="list-style-type: none"> - Expensive and slow to predict new instances - Must define a meaningful distance function - Performs poorly on high-dimensionality datasets | <ul style="list-style-type: none"> - Low-dimensional datasets - Computer security: intrusion detection - Fault detection in semiconductor manufacturing - Video content retrieval - Gene expression - Protein-protein interaction |